

Лабораторная работа №3. Реализация сверточной нейронной сети

In [1]:

```
from __future__ import absolute_import, division, print_function, unicode_literals

# TensorFlow и tf.keras
# !python3 -m pip install keras
import tensorflow as tf
from tensorflow import keras
from keras import regularizers
from keras.layers import Dense, Conv2D, Flatten

# Вспомогательные библиотеки
import numpy as np
import matplotlib.pyplot as plt
import pdb
from six.moves import cPickle as pickle
import os
from scipy import ndimage
```

Using TensorFlow backend.

In [2]:

```
def extract_dataset():
    with open('../data/notMNIST_sanit.pickle', 'rb') as f:
        data = pickle.load(f)

    # reshape dataset because of error:
    # ValueError: Error when checking input: expected conv2d_input to have 4 dimensions
    # but got array with shape (200000, 28, 28)
    for key, dataset in data.items():
        data[key] = dataset.reshape(*dataset.shape, 1)
    return data
```

In [3]:

```
dataset = extract_dataset()
train_images = dataset['train_dataset']
train_labels = dataset['train_labels']
valid_images = dataset['valid_dataset']
valid_labels = dataset['valid_labels']
test_images = dataset['test_dataset']
test_labels = dataset['test_labels']
```

Задание 1. Реализуйте нейронную сеть с двумя сверточными слоями, и одним полносвязным с нейронами с кусочно-линейной функцией активации. Какова точность построенной модели?

In [4]:

```

# kernel_size=3 – размер ядра 3x3.
# Функция активации 'relu' ( Rectified Linear Activation ),
# 64 это число ядер свертки( сколько признаков будем искать).
# Flatten() – слой, преобразующий 2D-данные в 1D-данные.
conv2d_model = keras.Sequential([
    keras.layers.Conv2D(8, kernel_size=3, activation='relu', input_shape=(28, 28, 1),
    keras.layers.Conv2D(16, kernel_size=3, activation='relu', input_shape=(28, 28, 1),
    keras.layers.Flatten(),
    keras.layers.Dense(100, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

conv2d_model.compile(optimizer='sgd',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

conv2d_model.summary()

conv2d_history = conv2d_model.fit(train_images,
                                train_labels,
                                epochs=10,
                                validation_data=(valid_images, valid_labels))

test_loss, test_acc = conv2d_model.evaluate(test_images, test_labels, verbose=2)

print('\nТочность на проверочных данных:', test_acc)

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 8)	80
conv2d_1 (Conv2D)	(None, 24, 24, 16)	1168
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 100)	921700
dense_1 (Dense)	(None, 10)	1010

```

Total params: 923,958
Trainable params: 923,958
Non-trainable params: 0

```

Train on 200000 samples, validate on 16911 samples

Epoch 1/10

```

200000/200000 [=====] - 79s 395us/sample - loss: 0.5508 - accuracy: 0.8404 - val_loss: 0.4634 - val_accuracy: 0.8659

```

Epoch 2/10

```

200000/200000 [=====] - 81s 406us/sample - loss: 0.3841 - accuracy: 0.8866 - val_loss: 0.4068 - val_accuracy: 0.8851

```

Epoch 3/10

```

200000/200000 [=====] - 81s 403us/sample - loss: 0.3310 - accuracy: 0.9018 - val_loss: 0.3706 - val_accuracy: 0.8957

```

Epoch 4/10

```
200000/200000 [=====] - 81s 404us/sample - loss: 0.2978 - accuracy: 0.9114 - val_loss: 0.3543 - val_accuracy: 0.8986
Epoch 5/10
200000/200000 [=====] - 81s 405us/sample - loss: 0.2725 - accuracy: 0.9183 - val_loss: 0.3469 - val_accuracy: 0.9012
Epoch 6/10
200000/200000 [=====] - 81s 405us/sample - loss: 0.2511 - accuracy: 0.9244 - val_loss: 0.3455 - val_accuracy: 0.9045
Epoch 7/10
200000/200000 [=====] - 81s 405us/sample - loss: 0.2317 - accuracy: 0.9302 - val_loss: 0.3467 - val_accuracy: 0.9008
Epoch 8/10
200000/200000 [=====] - 81s 405us/sample - loss: 0.2134 - accuracy: 0.9354 - val_loss: 0.3429 - val_accuracy: 0.9043
Epoch 9/10
200000/200000 [=====] - 81s 404us/sample - loss: 0.1960 - accuracy: 0.9405 - val_loss: 0.3523 - val_accuracy: 0.9034
Epoch 10/10
200000/200000 [=====] - 81s 407us/sample - loss: 0.1791 - accuracy: 0.9458 - val_loss: 0.3586 - val_accuracy: 0.9036
8722/8722 - 1s - loss: 0.1543 - accuracy: 0.9559
```

Точность на проверочных данных: 0.95585877

Задание 2. Замените один из сверточных слоев на слой, реализующий операцию пулинга (Pooling) с функцией максимума или среднего. Как это повлияло на точность классификатора?

Helpful links: <https://neurohive.io/ru/osnovy-data-science/glubokaya-svertochnaja-nejronnaja-set/>
(<https://neurohive.io/ru/osnovy-data-science/glubokaya-svertochnaja-nejronnaja-set/>)
<https://medium.com/@congyuzhou/pooling-%D1%81%D0%BB%D0%BE%D0%B9-dbe00ef48eab>
(<https://medium.com/@congyuzhou/pooling-%D1%81%D0%BB%D0%BE%D0%B9-dbe00ef48eab>)

In [5]:

```

maxpool_model = keras.Sequential([
    keras.layers.Conv2D(8, kernel_size=3, activation='relu', input_shape=(28, 28, 1),
    keras.layers.MaxPooling2D(),
    keras.layers.Flatten(),
    keras.layers.Dense(100, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

maxpool_model.compile(optimizer='sgd',
                      loss='sparse_categorical_crossentropy',
                      metrics=['accuracy'])

maxpool_model.summary()

maxpool_history = maxpool_model.fit(train_images,
                                    train_labels,
                                    epochs=10,
                                    validation_data=(valid_images, valid_labels))

test_loss, test_acc = maxpool_model.evaluate(test_images, test_labels, verbose=2)

print('\nТочность на проверочных данных:', test_acc)

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d (MaxPooling2D)	(None, 13, 13, 8)	0
flatten_1 (Flatten)	(None, 1352)	0
dense_2 (Dense)	(None, 100)	135300
dense_3 (Dense)	(None, 10)	1010
Total params: 136,390		
Trainable params: 136,390		
Non-trainable params: 0		

Train on 200000 samples, validate on 16911 samples

Epoch 1/10

200000/200000 [=====] - 36s 178us/sample - loss: 0.6299 - accuracy: 0.8208 - val_loss: 0.5312 - val_accuracy: 0.8454

Epoch 2/10

200000/200000 [=====] - 35s 177us/sample - loss: 0.4515 - accuracy: 0.8668 - val_loss: 0.4581 - val_accuracy: 0.8662

Epoch 3/10

200000/200000 [=====] - 36s 178us/sample - loss: 0.4034 - accuracy: 0.8808 - val_loss: 0.4335 - val_accuracy: 0.8754

Epoch 4/10

200000/200000 [=====] - 36s 178us/sample - loss: 0.3748 - accuracy: 0.8895 - val_loss: 0.4133 - val_accuracy: 0.8823

Epoch 5/10

```
200000/200000 [=====] - 36s 178us/sample - loss: 0.3542 - accuracy: 0.8961 - val_loss: 0.3980 - val_accuracy: 0.8856
Epoch 6/10
200000/200000 [=====] - 36s 178us/sample - loss: 0.3383 - accuracy: 0.9007 - val_loss: 0.3845 - val_accuracy: 0.8894
Epoch 7/10
200000/200000 [=====] - 36s 179us/sample - loss: 0.3249 - accuracy: 0.9039 - val_loss: 0.3844 - val_accuracy: 0.8888
Epoch 8/10
200000/200000 [=====] - 35s 177us/sample - loss: 0.3133 - accuracy: 0.9075 - val_loss: 0.3717 - val_accuracy: 0.8922
Epoch 9/10
200000/200000 [=====] - 36s 178us/sample - loss: 0.3033 - accuracy: 0.9103 - val_loss: 0.3718 - val_accuracy: 0.8920
Epoch 10/10
200000/200000 [=====] - 36s 178us/sample - loss: 0.2941 - accuracy: 0.9129 - val_loss: 0.3657 - val_accuracy: 0.8960
8722/8722 - 1s - loss: 0.1731 - accuracy: 0.9493
```

Точность на проверочных данных: 0.94932353

Задание 3. Реализуйте классическую архитектуру сверточных сетей LeNet-5

(<http://yann.lecun.com/exdb/lenet/>) (<http://yann.lecun.com/exdb/lenet/>).

padding: <https://medium.com/@congyuzhou/padding-32266fa95816>

(<https://medium.com/@congyuzhou/padding-32266fa95816>) <https://medium.com/@congyuzhou/lenet-5-%D1%81%D0%B2%D0%BE%D0%B8%D0%BC%D0%B8-%D1%80%D1%83%D0%BA%D0%B0%D0%BC%D0%B8-b60ae3727cd3>
(<https://medium.com/@congyuzhou/lenet-5-%D1%81%D0%B2%D0%BE%D0%B8%D0%BC%D0%B8-%D1%80%D1%83%D0%BA%D0%B0%D0%BC%D0%B8-b60ae3727cd3>)

In [6]:

```
# stride - war
lenet5_model = keras.Sequential([
    keras.layers.Conv2D(filters=6, kernel_size=[5,5], padding='same', activation='relu'),
    keras.layers.MaxPooling2D(pool_size=[2,2], strides=2),
    keras.layers.Conv2D(filters=16, kernel_size=[5,5], padding='valid', activation='relu'),
    keras.layers.MaxPooling2D(pool_size=[2,2], strides=2),
    keras.layers.Flatten(),
    keras.layers.Dense(units=120, activation='relu'),
    keras.layers.Dense(units=84, activation='relu'),
    keras.layers.Dense(units=10, activation='softmax')
])

lenet5_model.compile(optimizer='sgd',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

lenet5_history = lenet5_model.fit(train_images,
                                train_labels,
                                epochs=10,
                                validation_data=(valid_images, valid_labels))

lenet5_model.summary()

test_loss, test_acc = lenet5_model.evaluate(test_images, test_labels, verbose=2)

print('\nТочность на проверочных данных:', test_acc)
```

Train on 200000 samples, validate on 16911 samples

Epoch 1/10

200000/200000 [=====] - 67s 336us/sample - loss: 0.6139 - accuracy: 0.8175 - val_loss: 0.4642 - val_accuracy: 0.8650

Epoch 2/10

200000/200000 [=====] - 67s 336us/sample - loss: 0.3927 - accuracy: 0.8833 - val_loss: 0.4094 - val_accuracy: 0.8798

Epoch 3/10

200000/200000 [=====] - 68s 339us/sample - loss: 0.3551 - accuracy: 0.8929 - val_loss: 0.3850 - val_accuracy: 0.8856

Epoch 4/10

200000/200000 [=====] - 67s 337us/sample - loss: 0.3323 - accuracy: 0.8997 - val_loss: 0.3659 - val_accuracy: 0.8921

Epoch 5/10

200000/200000 [=====] - 68s 338us/sample - loss: 0.3164 - accuracy: 0.9040 - val_loss: 0.3538 - val_accuracy: 0.8955

Epoch 6/10

200000/200000 [=====] - 68s 338us/sample - loss: 0.3040 - accuracy: 0.9075 - val_loss: 0.3497 - val_accuracy: 0.8966

Epoch 7/10

200000/200000 [=====] - 68s 338us/sample - loss: 0.2942 - accuracy: 0.9105 - val_loss: 0.3354 - val_accuracy: 0.9005

Epoch 8/10

200000/200000 [=====] - 68s 338us/sample - loss: 0.2848 - accuracy: 0.9132 - val_loss: 0.3271 - val_accuracy: 0.9010

Epoch 9/10

200000/200000 [=====] - 68s 342us/sample - loss: 0.2772 - accuracy: 0.9145 - val_loss: 0.3237 - val_accuracy: 0.9030

Epoch 10/10

200000/200000 [=====] - 68s 340us/sample - loss: 0.2696 - accuracy: 0.9169 - val_loss: 0.3192 - val_accuracy: 0.9047

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	multiple	156
max_pooling2d_1 (MaxPooling2D)	multiple	0
conv2d_4 (Conv2D)	multiple	2416
max_pooling2d_2 (MaxPooling2D)	multiple	0
flatten_2 (Flatten)	multiple	0
dense_4 (Dense)	multiple	48120
dense_5 (Dense)	multiple	10164
dense_6 (Dense)	multiple	850
=====		
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		

8722/8722 - 1s - loss: 0.1430 - accuracy: 0.9539

Точность на проверочных данных: 0.95390964

In [7]:

```
def plot_history(histories, key='binary_crossentropy'):
    plt.figure(figsize=(16,10))


    for name, history in histories:
        val = plt.plot(history.epoch, history.history['val_' + key],
                        '--', label=name.title()+' Val')
        plt.plot(history.epoch, history.history[key], color=val[0].get_color(),
                 label=name.title()+' Train')

    plt.xlabel('Epochs')
    plt.ylabel(key.replace('_', ' ').title())
    plt.legend()

    plt.xlim([0,max(history.epoch)])

    plt.show()

plot_history(
    [
        ('conv2d', conv2d_history),
        ('maxpool', maxpool_history),
        ('lenet5', lenet5_history)
    ],
    key='loss')
```



In []: