

Лабораторная работа №4. Реализация приложения по распознаванию номеров домов.

In [3]:

```
# TensorFlow и tf.keras
import tensorflow as tf
from tensorflow import keras
from keras import regularizers

import numpy as np
import matplotlib.pyplot as plt
import pdb
import os
import scipy.io
from sklearn.model_selection import train_test_split
import tarfile
from six.moves import cPickle as pickle

l2_regularization = 1e-4
```

Using TensorFlow backend.

Задание 1. Реализуйте глубокую нейронную сеть (полносвязную или сверточную) и обучите ее на синтетических данных (например, наборы MNIST (<http://yann.lecun.com/exdb/mnist/>) или notMNIST).

In [4]:

```
# Extract data

def extract_dataset():
    with open('../data/notMNIST_sanit.pickle', 'rb') as f:
        data = pickle.load(f)

    # reshape dataset because of error:
    # ValueError: Error when checking input: expected conv2d_input to have 4 dimensions
    # but got array with shape (2000000, 28, 28)
    for key, dataset in data.items():
        data[key] = dataset.reshape(*dataset.shape, 1)
    return data

def image_name(index):
    return chr(ord('A') + index)

dataset = extract_dataset()
train_images = dataset['train_dataset']
train_labels = dataset['train_labels']
valid_images = dataset['valid_dataset']
valid_labels = dataset['valid_labels']
test_images = dataset['test_dataset']
test_labels = dataset['test_labels']
```

In [5]:

```
numbers_model = keras.Sequential([
    keras.layers.Conv2D(filters=16, kernel_size=[5,5], padding='same', activation='relu'),
    keras.layers.MaxPooling2D(pool_size=[2,2], strides=2),
    keras.layers.Dropout(0.5),
    keras.layers.Conv2D(filters=32, kernel_size=[5,5], padding='valid', activation='relu'),
    keras.layers.MaxPooling2D(pool_size=[2,2], strides=2),
    keras.layers.Dropout(0.5),
    keras.layers.Flatten(),
    keras.layers.Dense(units=120, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(units=84, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(units=10, activation = 'softmax')
])

optimizer = keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-08)
numbers_model.compile(optimizer=optimizer,
                      loss='sparse_categorical_crossentropy',
                      metrics=['accuracy'])

numbers_model.summary()

numbers_model_history = numbers_model.fit(train_images,
                                         train_labels,
                                         epochs=10,
                                         validation_data=(valid_images, valid_labels))

test_loss, test_acc = numbers_model.evaluate(test_images, test_labels, verbose=2)

print('\nТочность на проверочных данных:', test_acc)
```

Задание 2. После уточнения модели на синтетических данных попробуйте обучить ее на реальных данных (набор Google Street View). Что изменилось в модели?

In [6]:

```
# Extracting data
```

```
dataset_paths = [('test', "../data/housenames/test.tar.gz"),  
                  ('train', "../data/housenames/train.tar.gz")]
```

```
for name, path in dataset_paths:  
    if not os.path.exists('../data/housenames/' + name):  
        print(name)  
        tf = tarfile.open(path)  
        files = tf.extractall('../data/housenames')  
        tf.close()
```

```
houses_train_dataset = scipy.io.loadmat(os.path.join('../data/housenames/', 'train_32'))  
houses_test_dataset = scipy.io.loadmat(os.path.join('../data/housenames/', 'test_32'))
```

```
# Formatting of data
```

```
# 1) move axis to be like this: (num of images, h, w, dimensions)
```

```
houses_train_dataset['X'] = np.moveaxis(houses_train_dataset['X'], 3, 0)  
houses_test_dataset['X'] = np.moveaxis(houses_test_dataset['X'], 3, 0)
```

```
# 2) subtract 1 from labels because they start from 1
```

```
houses_train_dataset['y'] = houses_train_dataset['y'] - 1  
houses_test_dataset['y'] = houses_test_dataset['y'] - 1
```

```
houses_train_images, houses_valid_images, houses_train_labels, houses_valid_labels =  
    houses_train_dataset['X'],  
    houses_train_dataset['y']  
)
```

```
houses_test_images = houses_test_dataset['X']  
houses_test_labels = houses_test_dataset['y']
```

In [7]:

```

houses_model_v1 = keras.Sequential([
    keras.layers.Conv2D(32, 3, padding='same', input_shape=(32, 32, 3), activation='relu'),
    keras.layers.MaxPooling2D(pool_size=(2, 2)),
    keras.layers.Conv2D(64, 3, padding='same', activation='relu'),
    keras.layers.MaxPooling2D(pool_size=(2, 2)),
    keras.layers.Flatten(),
    keras.layers.Dense(512, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

houses_model_v1.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
houses_model_v1.summary()
houses_model_v1_history = houses_model_v1.fit(houses_train_images,
                                              houses_train_labels,
                                              epochs=10,
                                              validation_data=(houses_valid_images,
                                                                houses_valid_labels),
                                              batch_size=128)

test_loss, test_acc = houses_model_v1.evaluate(houses_test_images, houses_test_labels)

print('\nТочность на проверочных данных:', test_acc)

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 512)	2097664
dense_1 (Dense)	(None, 10)	5130
=====		

Total params: 2,122,186  
Trainable params: 2,122,186  
Non-trainable params: 0

Train on 54942 samples, validate on 18315 samples

```

Epoch 1/10
54942/54942 [=====] - 53s 960us/sample - loss: 4.2344 - accuracy: 0.5337 - val_loss: 0.9970 - val_accuracy: 0.6888
Epoch 2/10
54942/54942 [=====] - 57s 1ms/sample - loss: 0.7778 - accuracy: 0.7656 - val_loss: 0.7380 - val_accuracy: 0.7844
Epoch 3/10
54942/54942 [=====] - 53s 961us/sample - loss: 0.6216 - accuracy: 0.8141 - val_loss: 0.6754 - val_accuracy: 0.8044
Epoch 4/10
54942/54942 [=====] - 54s 976us/sample - loss: 0.5263 - accuracy: 0.8429 - val_loss: 0.6861 - val_accuracy: 0.8029
Epoch 5/10
54942/54942 [=====] - 58s 1ms/sample - loss: 0.4714 - accuracy: 0.8571 - val_loss: 0.6655 - val_accuracy: 0.8140

```

```
Epoch 6/10
54942/54942 [=====] - 56s 1ms/sample - loss:
0.4254 - accuracy: 0.8702 - val_loss: 0.6651 - val_accuracy: 0.8197
Epoch 7/10
54942/54942 [=====] - 63s 1ms/sample - loss:
0.3788 - accuracy: 0.8816 - val_loss: 0.6942 - val_accuracy: 0.8194
Epoch 8/10
54942/54942 [=====] - 70s 1ms/sample - loss:
0.3320 - accuracy: 0.8964 - val_loss: 0.6925 - val_accuracy: 0.8239
Epoch 9/10
54942/54942 [=====] - 63s 1ms/sample - loss:
0.2998 - accuracy: 0.9062 - val_loss: 0.7181 - val_accuracy: 0.8223
Epoch 10/10
54942/54942 [=====] - 55s 1ms/sample - loss:
0.2718 - accuracy: 0.9145 - val_loss: 0.7385 - val_accuracy: 0.8216
26032/26032 - 8s - loss: 0.9590 - accuracy: 0.7974
```

Точность на проверочных данных: 0.7974032

In [14]:

```
# version 2: + regularization to dense layer, remove batch_size during fitting + 1

houses_model_v2 = keras.Sequential([
    keras.layers.Conv2D(32, 3, activation='relu', padding='same', input_shape=(32,
    keras.layers.MaxPooling2D(pool_size=(2, 2)),

    keras.layers.Conv2D(64, 3, activation='relu', padding='same'),
    keras.layers.MaxPooling2D(pool_size=(2, 2)),

    keras.layers.Flatten(),
    keras.layers.Dense(units=512, activation='relu', kernel_regularizer=regularizer
    keras.layers.Dropout(0.25),

    keras.layers.Dense(units=10, activation = 'softmax')
])

houses_model_v2.compile(optimizer='adam',
                        loss='sparse_categorical_crossentropy',
                        metrics=['accuracy'])

houses_model_v2.summary()
houses_model_v2_history = houses_model_v2.fit(houses_train_images,
                                              houses_train_labels,
                                              epochs=10,
                                              validation_data=(houses_valid_images,
test_loss, test_acc = houses_model_v2.evaluate(houses_test_images, houses_test_lab

print('\nТочность на проверочных данных:', test_acc)
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_11 (MaxPooling)	(None, 16, 16, 32)	0
conv2d_12 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_12 (MaxPooling)	(None, 8, 8, 64)	0
flatten_5 (Flatten)	(None, 4096)	0
dense_8 (Dense)	(None, 512)	2097664
dropout_9 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 10)	5130
Total params: 2,122,186		
Trainable params: 2,122,186		
Non-trainable params: 0		

Train on 54942 samples, validate on 18315 samples

Epoch 1/10

54942/54942 [=====] - 78s 1ms/sample - loss:

1.5694 - accuracy: 0.6620 - val\_loss: 0.7330 - val\_accuracy: 0.8050

Epoch 2/10

54942/54942 [=====] - 91s 2ms/sample - loss:

```
0.6932 - accuracy: 0.8170 - val_loss: 0.7413 - val_accuracy: 0.8068
Epoch 3/10
54942/54942 [=====] - 113s 2ms/sample - loss:
0.6227 - accuracy: 0.8384 - val_loss: 0.5964 - val_accuracy: 0.8548
Epoch 4/10
54942/54942 [=====] - 119s 2ms/sample - loss:
0.5733 - accuracy: 0.8532 - val_loss: 0.6003 - val_accuracy: 0.8555
Epoch 5/10
54942/54942 [=====] - 118s 2ms/sample - loss:
0.5365 - accuracy: 0.8673 - val_loss: 0.6095 - val_accuracy: 0.8577
Epoch 6/10
54942/54942 [=====] - 119s 2ms/sample - loss:
0.5047 - accuracy: 0.8765 - val_loss: 0.5797 - val_accuracy: 0.8664
Epoch 7/10
54942/54942 [=====] - 117s 2ms/sample - loss:
0.4924 - accuracy: 0.8814 - val_loss: 0.5923 - val_accuracy: 0.8682
Epoch 8/10
54942/54942 [=====] - 120s 2ms/sample - loss:
0.4760 - accuracy: 0.8899 - val_loss: 0.6448 - val_accuracy: 0.8554
Epoch 9/10
54942/54942 [=====] - 118s 2ms/sample - loss:
0.4616 - accuracy: 0.8959 - val_loss: 0.6571 - val_accuracy: 0.8638
Epoch 10/10
54942/54942 [=====] - 117s 2ms/sample - loss:
0.4622 - accuracy: 0.8975 - val_loss: 0.6346 - val_accuracy: 0.8678
26032/26032 - 14s - loss: 0.7329 - accuracy: 0.8552
```

Точность на проверочных данных: 0.8552167

In [15]:

```
def plot_history(histories, key='binary_crossentropy'):
    plt.figure(figsize=(16,10))


    for name, history in histories:
        val = plt.plot(history.epoch, history.history['val_' + key],
                        '--', label=name.title()+' Val')
        plt.plot(history.epoch, history.history[key], color=val[0].get_color(),
                 label=name.title()+' Train')

    plt.xlabel('Epochs')
    plt.ylabel(key.replace('_', ' ').title())
    plt.legend()

    plt.xlim([0,max(history.epoch)])

    plt.show()

plot_history(
[
    ('houses_v1', houses_model_v1_history),
    ('houses_v2', houses_model_v2_history),
], key='loss')
```



In [ ]: