

Лабораторная работа №7. Рекуррентные нейронные сети для анализа текста

In [2]:

```
# TensorFlow и tf.keras
import tensorflow as tf
from tensorflow import keras
from keras import regularizers

import numpy as np
import matplotlib.pyplot as plt
import pdb
import os
import scipy.io
from sklearn.model_selection import train_test_split
import tarfile
from six.moves import cPickle as pickle
import zipfile
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import re
from collections import defaultdict
```

Using TensorFlow backend.

Задание 1. Загрузите данные. Преобразуйте текстовые файлы во внутренние структуры данных, которые используют индексы вместо слов.

In []:

```
# загрузка и форматирование данных

# used only once
# !bash ../data/movie_reviews/preprocess_reviews.sh
```

Скрипт из открытого доступа (<https://gist.github.com/aaronkub/09985a47740bda278712e1dd78e482cf>) объединяет все негативные рецензии и все позитивные в одни файлы. Он исправлен, чтобы файл был не один.

In [3]:

```
dataset_path = '../data/movie_reviews/aclImdb/data'

def extract_data(filename):
    dataset = []
    with open(os.path.join(dataset_path, filename)) as f:
        for lines in f.readlines():
            dataset.append(lines.strip())
    return dataset

train_pos_reviews = extract_data('full_pos_train.txt')
train_neg_reviews = extract_data('full_neg_train.txt')
test_pos_reviews = extract_data('full_pos_test.txt')
test_neg_reviews = extract_data('full_neg_test.txt')

train_labels = np.concatenate((np.ones(len(train_pos_reviews)), np.zeros(len(train_neg_reviews))))
test_labels = np.concatenate((np.ones(len(test_pos_reviews)), np.zeros(len(test_neg_reviews))))
```

In [17]:

```
print('Train positive lines: ', len(train_pos_reviews))
print('Train negative lines: ', len(train_neg_reviews))
print('Train labels: ', len(train_labels))
print('Test positive lines: ', len(test_pos_reviews))
print('Test negative lines: ', len(test_neg_reviews))
print('Test labels: ', len(test_labels))
```

```
Train positive lines: 12500
Train negative lines: 12500
Train labels: 25000
Test positive lines: 12500
Test negative lines: 12500
Test labels: 25000
```

<https://towardsdatascience.com/sentiment-analysis-with-python-part-1-5ce197074184>
[\(https://towardsdatascience.com/sentiment-analysis-with-python-part-1-5ce197074184\)](https://towardsdatascience.com/sentiment-analysis-with-python-part-1-5ce197074184)

Препроцессим данные и делаем их чище (убираем заглавные, знаки препинания и так далее).

In [4]:

```
REPLACE_NO_SPACE = re.compile("[.;:!\"'?,\\\"()\\[\\]]")
REPLACE_WITH_SPACE = re.compile("<br\\s*/><br\\s*/>|\\(\\-)|\\(\\/\\)")

def preprocess_reviews(reviews):
    reviews = [REPLACE_NO_SPACE.sub("", line.lower()) for line in reviews]
    reviews = [REPLACE_WITH_SPACE.sub(" ", line) for line in reviews]

    return reviews

train_pos_reviews = preprocess_reviews(train_pos_reviews)
train_neg_reviews = preprocess_reviews(train_neg_reviews)
test_pos_reviews = preprocess_reviews(test_pos_reviews)
test_neg_reviews = preprocess_reviews(test_neg_reviews)
```

In [5]:

```
# https://stackoverflow.com/questions/51956000/what-does-keras-tokenizer-method-exa
# https://stackoverflow.com/questions/42943291/what-does-keras-io-preprocessing-seq

tokenizer = keras.preprocessing.text.Tokenizer(num_words=10000, split=' ')
tokenizer.fit_on_texts(train_pos_reviews + train_neg_reviews)

X = tokenizer.texts_to_sequences(train_pos_reviews + train_neg_reviews + test_pos_r
X = tf.keras.preprocessing.sequence.pad_sequences(X)

X_train, X_test = np.split(X, [25000])

y_train = train_labels
y_test = test_labels

X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train, test_size =
```

In [40]:

```
print('Train X: ', X_train.shape)
print('Valid X: ', X_valid.shape)
print('Test X: ', X_test.shape)
print('Train y: ', y_train.shape)
print('Valid y: ', y_valid.shape)
print('Test y: ', y_test.shape)
```

```
Train X: (18750, 2176)
Valid X: (6250, 2176)
Test X: (25000, 2176)
Train y: (18750,)
Valid y: (6250,)
Test y: (25000,)
```

Задание 2. Реализуйте и обучите двунаправленную рекуррентную сеть (LSTM или GRU). Какого качества классификации удалось достичь?

In [41]:

```
embedding_vector_length = 32

lstm_model = keras.models.Sequential([
    keras.layers.Embedding(10000, embedding_vector_length, input_length=X_train.shape[1]),
    keras.layers.LSTM(100),
    keras.layers.Dense(1, activation='sigmoid')
])

lstm_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 2176, 32)	320000
lstm_1 (LSTM)	(None, 100)	53200
dense_1 (Dense)	(None, 1)	101
Total params: 373,301		
Trainable params: 373,301		
Non-trainable params: 0		

In [45]:

```
lstm_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
lstm_model_history = lstm_model.fit(X_train, y_train, validation_data=(X_valid, y_valid),
```

Train on 18750 samples, validate on 6250 samples

Epoch 1/3

18750/18750 [=====] - 945s 50ms/sample - loss: 0.4725 - accuracy: 0.7867 - val_loss: 0.3591 - val_accuracy: 0.8494

Epoch 2/3

18750/18750 [=====] - 947s 50ms/sample - loss: 0.2834 - accuracy: 0.8877 - val_loss: 0.3431 - val_accuracy: 0.8530

Epoch 3/3

18750/18750 [=====] - 909s 48ms/sample - loss: 0.2064 - accuracy: 0.9218 - val_loss: 0.4535 - val_accuracy: 0.7906

In [46]:

```
test_loss, test_acc = lstm_model.evaluate(X_test, y_test, verbose=2)

print('\nТочность на проверочных данных:', test_acc)
```

25000/25000 - 310s - loss: 0.4565 - accuracy: 0.7914

Точность на проверочных данных: 0.7914


In [47]:

```
def plot_history(histories, key='binary_crossentropy'):
    plt.figure(figsize=(16,10))

    for name, history in histories:
        val = plt.plot(history.epoch, history.history['val_' + key],
                        '--', label=name.title()+' Val')
        plt.plot(history.epoch, history.history[key], color=val[0].get_color(),
                 label=name.title()+' Train')

    plt.xlabel('Epochs')
    plt.ylabel(key.replace('_', ' ').title())
    plt.legend()
    plt.xlim([0,max(history.epoch)])
    plt.show()

plot_history([('lstm model', lstm_model_history)], key='loss')
```



Задание 3. Используйте индексы слов и их различное внутреннее представление (word2vec, glove). Как влияет данное преобразование на качество классификации?

In [6]:

```
def load_embeddings(filename):
    embeddings_index = {}

    with open(filename, 'r') as in_file:
        for line in in_file:
            values = line.split()

            try:
                word = values[0]
                embeddings_index[word] = np.asarray(values[1:], dtype=np.float32)
            except:
                pass

    return embeddings_index

def get_embedding_matrix(embedding, vocab):
    vocab_size = len(vocab) + 1
    weight_matrix = np.zeros((vocab_size, 300))

    for word, i in vocab.items():
        weight_matrix[i] = embedding.get(word)

    return weight_matrix

embedding = load_embeddings("../data/movie_reviews/glove/glove.840B.300d.txt")
weight_matrix = get_embedding_matrix(embedding, tokenizer.word_index)
```

In [9]:

```
vocab_size = len(tokenizer.word_index) + 1

glove_model = keras.models.Sequential([
    keras.layers.Embedding(vocab_size, 300, weights=[weight_matrix], trainable=False),
    keras.layers.LSTM(100),
    keras.layers.Dense(1, activation='sigmoid')
])

glove_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, None, 300)	28131600
=====		
lstm_1 (LSTM)	(None, 100)	160400
=====		
dense_1 (Dense)	(None, 1)	101
=====		
Total params: 28,292,101		
Trainable params: 160,501		
Non-trainable params: 28,131,600		

In []:

```
glove_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
glove_model_history = glove_model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10)
```

In []:

```
test_loss, test_acc = glove_model.evaluate(X_test, y_test, verbose=2)
print('\nТочность на проверочных данных:', test_acc)
```

In []:

```
plot_history([('glove model', glove_model_history)], key='loss')
```

Задание 4.

Поэкспериментируйте со структурой сети (добавьте больше рекуррентных, полносвязных или сверточных слоев). Как это повлияло на качество классификации?

In []:

```
embedding_vector_length = 32

diff_model = keras.models.Sequential([
    keras.layers.Embedding(10000, embedding_vector_length, input_length=X_train.shape[1]),
    keras.layers.Conv1D(filters=32, kernel_size=3, padding='same', activation='relu'),
    keras.layers.MaxPooling1D(pool_size=2),
    keras.layers.Dropout(0.2),
    keras.layers.LSTM(100, dropout=0.2, recurrent_dropout=0.2),
    keras.layers.Dense(1, activation='sigmoid')
])

diff_model.summary()
```

In []:

```
diff_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
diff_model_history = diff_model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10)
```

In []:

```
test_loss, test_acc = diff_model.evaluate(X_test, y_test, verbose=2)
print('\nТочность на проверочных данных:', test_acc)
```

In []:

```
plot_history([('diff model', diff_model_history)], key='loss')
```